

Sai Satwik Bikumandla

Albany, New York, United States | bikumandlasaisatwik@gmail.com | (518) 614-1904 | [GitHub](#) | [LinkedIn](#) | [Website](#)

PROFESSIONAL SUMMARY

Software Engineer with full-stack knowledge who builds production grade software applications with AI/ML capabilities.

EDUCATION

University at Albany, State University of New York Master of Science, <i>Computer Science</i>	Aug 2024 - May 2026
Sri Indu College of Engineering & Technology (JNTUH) Bachelor of Technology, <i>Computer Science & Information Technology</i>	Aug 2020 - Apr 2024

SKILLS

Programming Languages: Python, JavaScript, TypeScript, Java, C++
Frameworks & Libraries: React.js, Next.js, Node.js, Flask, TensorFlow/Keras, YOLO, LSTM/RNN
Databases: MySQL, MongoDB, Google Cloud SQL
Tools & Software: Git, LaTeX, Docker, MATLAB, Gemini API, REST APIs, NextAuth, JWT, HTML/CSS
Cloud Platforms: Google Cloud Platform
Soft Skills: Technical Documentation, Knowledge Transfer, Communication
Spoken Languages: Telugu (Native), Hindi (Fluent), English (Fluent), German (Actively Learning)

PROJECTS

LaTeX Resume Builder - Version-Controlled Resume Platform

- Designed and shipped a full-stack SaaS-style web application (Next.js, MongoDB, NextAuth) giving job seekers a GitHub-inspired resume workflow: multiple named repositories, commit snapshots, full version history, and rollback to any prior state - eliminating the problem of scattered, inconsistent resume files.
- Engineered a dual-render pipeline: an HTML live preview for zero-latency editing feedback alongside a server-side LaTeX compilation endpoint that exports polished, print-ready PDFs from any historical commit.
- Modeled resume versions as append-only, immutable snapshots (User → ResumeRepository → ResumeVersion), enforcing audit-trail integrity while keeping the UX approachable for non-technical users across 4 purpose-built pages.

Movie Recommendation System - Content-Based NLP Engine

- Built a content-based recommender over a dataset of 5,000+ films using TF-IDF vectorization and cosine similarity; applied a full NLP preprocessing pipeline (tokenization, lemmatization, stop-word removal) to improve feature signal quality.
- Evaluated recommendation relevance using Precision@K across K = 5, 10, and 20, enabling data-driven tuning of the similarity threshold and surfacing measurable lift over a random baseline.

Tennis Ball Detection & Tracking System

- Integrated Roboflow's pre-trained tennis ball detection model via REST API and built two production-quality tracking systems on top: a Basic Model with anti-jump filtering, exponential smoothing ($\alpha=0.85$), and multi-frame velocity extrapolation; and an Advanced Model adding court-side A/B classification, perspective-aware overlays, and configurable color schemes per court surface type.
- Implemented SORT (Simple Online and Realtime Tracking) for multi-object tracking with custom bounding-box conversion, per-track trajectory trails (50-frame memory), and frame-by-frame annotation overlays displaying Ball ID, detection count, and active track count.
- Built a TrackingAnalytics engine that generates structured JSON output per session - capturing velocity peaks, total trajectory distance, side-time distribution, detection success rate, and position maps - exported alongside processed video.
- Deployed as an interactive Gradio UI with model selection (Basic vs. Advanced), configurable start/end time segments, court color scheme selection, live analytics summary, session run history, and ZIP download of processed video + analytics JSON.

Next-Word Predictor - LSTM Language Model

- Trained an LSTM language model (Embedding → LSTM 150 units → Softmax Dense) in TensorFlow/Keras on 2,000 IMDB reviews; built the full data pipeline - tokenization, sliding-window sequencing (length 20), padding, and one-hot encoding - over a 20,000-token corpus with an 80/20 train-validation split.
- Implemented checkpoint-based model persistence: training runs once and subsequent sessions load saved weights instantly, enabling iterative experimentation without redundant computation.

Web Prototype Generator - AI-Assisted Code Generation Tool

- Led the Gemini API integration layer, building the prompt engineering pipeline that converts natural-language client briefs into functional, responsive HTML/CSS/JS prototypes with auto-embedded Unsplash imagery; backend service layer built with Flask.
- Introduced in-source AI-attribution tagging and documentation standards adopted across the team codebase, ensuring transparency on human vs. AI-generated contributions.

WORK EXPERIENCE

Rethink UX Website Design Intern	Oct 2022 - Dec 2022
- Designed and built client-facing website features in HTML, CSS, and JavaScript, taking assigned modules from requirements through delivery. - Implemented access-control logic to restrict page features by user role. - Wrote technical documentation for the features I built, making handoff easier for the team after the internship ended. - Participated in code reviews with senior engineers and incorporated feedback to meet production standards.	

StudyExperts Programming Content Intern	Jun 2022 - Aug 2022
- Created programming tutorials and technical knowledge-base articles covering Python, Java & data structures, serving as self-service resources for the platform's users. - Wrote, tested, and debugged code examples across multiple languages to ensure published content was accurate and runnable. - Demonstrated that strong written communication and deep product understanding are force multipliers for small engineering teams - a lesson that now directly informs how I document and advocate for my own projects.	